

The HTTP Protocol

RES, Lecture 4

Olivier Liechti



HAUTE ÉCOLE
D'INGÉNIERIE ET DE GESTION
DU CANTON DE VAUD

www.heig-vd.ch

Welcome to “the Web”



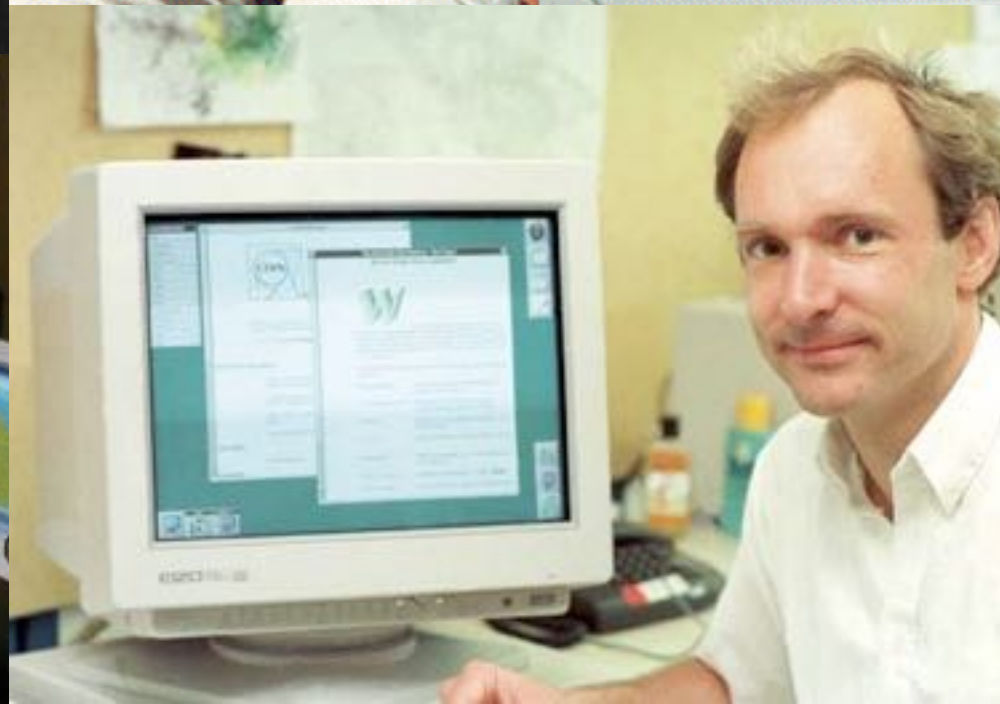
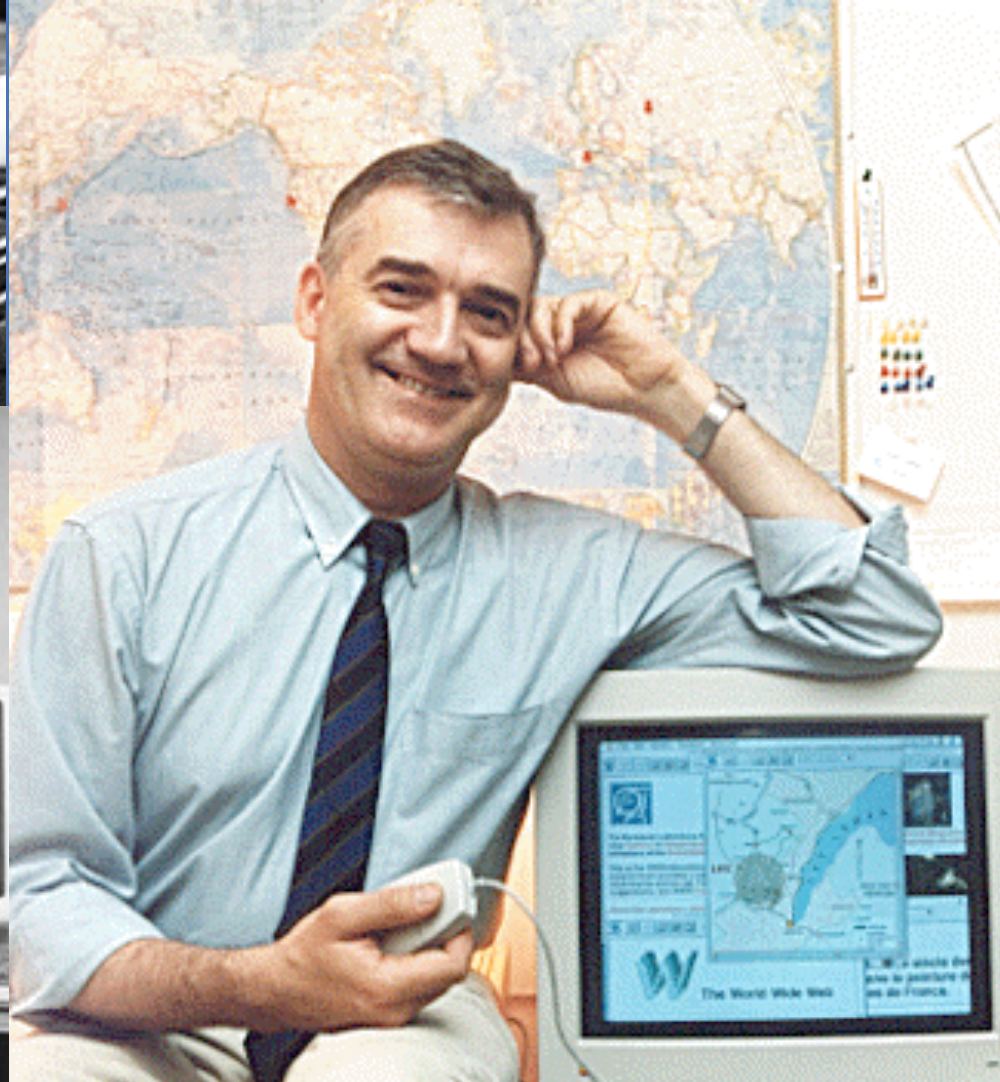
*It started as an (hypertext) **library**...*



*... which is now blending with the **physical** world*



*... which evolved to become an online application **platform**...*





http://en.wikipedia.org/wiki/Image:First_Web_Server.jpg



<http://www.w3.org/Consortium/technology>

What is HTTP?

The Hypertext Transfer Protocol (HTTP) is an application-level protocol for distributed, collaborative, **hypermedia information systems**.

It is a **generic, stateless**, protocol which can be used for many tasks beyond its use for hypertext, such as name servers and distributed object management systems, through **extension** of its request methods, error codes and headers [47].

A feature of HTTP is the typing and **negotiation of data representation**, allowing systems to be built independently of the data being transferred.

HTTP is one of the first two standards
of “the Web”



HTML

Markup Language to
create hypertext documents



HTTP

Protocol to **transfer** hypertext
documents (and other content)

What is HTTP?

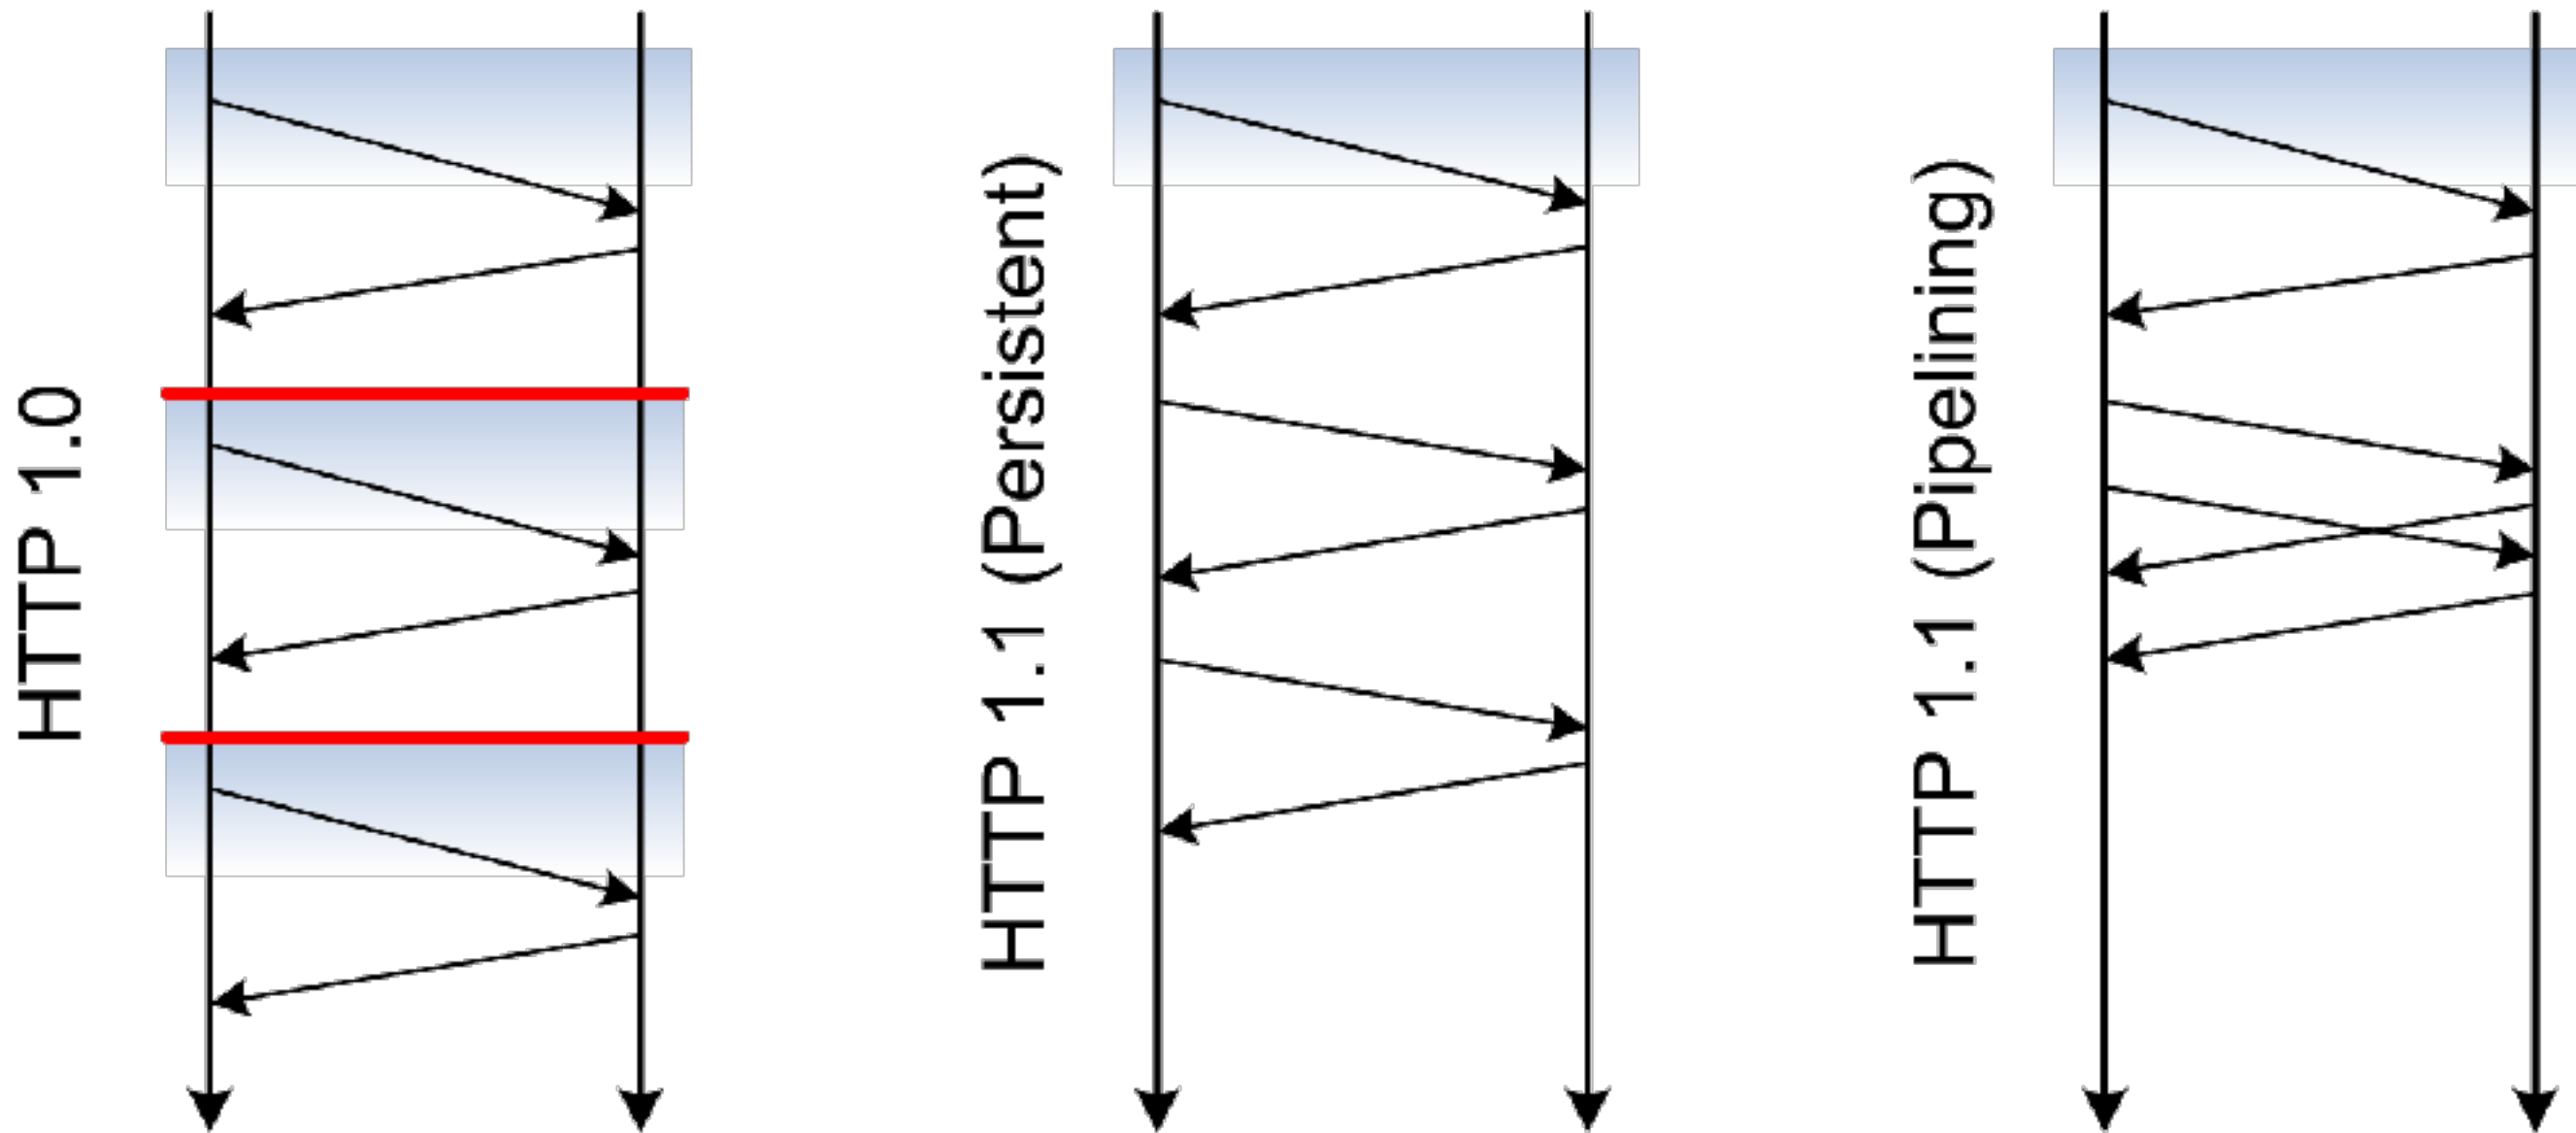
- **H**yper**T**ext **T**ransfer **P**rotocol.
- HTTP is an **application-level** protocol.
- HTTP is used to **transfer different types of payloads** (HTML, XML, JSON, PNG, MP4, WAV, etc.) between **clients** and **servers**
 - Sometimes, the client issues a request to GET (i.e. obtain, fetch, download) a payload from a server.
 - Sometimes, the client issues a request to POST (i.e. send, upload) a payload to a server.
- HTTP is built on top of **TCP** (the standard specifies that a server should accept requests on port 80).



The browser is an
HTTP client

The HTTP server accepts TCP connection
requests (by default on port 80)

HTTP & TCP Connections




[http://dret.net/lectures/web-fall07/foundations#\(20\)](http://dret.net/lectures/web-fall07/foundations#(20))


<http://www.apacheweek.com/features/http11>

Looking at a Conversation...





```
GET / HTTP/1.1 CRLF
Host: www.nodejs.org CRLF
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.8; rv:28.0) Gecko/20100101 Firefox/28.0 CRLF
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8 CRLF
Accept-Language: en-us,en;q=0.8,fr;q=0.5,fr-fr;q=0.3 CRLF
Accept-Encoding: gzip, deflate CRLF
Cookie: __utma=212211339.431073283.1392993818.1395308748.1395311696.27;
__utmz=212211339.1395311696.27.19.utmcsr=stackoverflow.com|utmccn=(referral)|utmcmd=referral
|utmcct=/questions/7776452/retrieving-a-list-of-network-interfaces-in-node-js- CRLF
Connection: keep-alive CRLF
CRLF
```



```
HTTP/1.1 200 OK CRLF
Server: nginx CRLF
Date: Sat, 05 Apr 2014 11:45:48 GMT CRLF
Content-Type: text/html CRLF
Content-Length: 6368 CRLF
Last-Modified: Tue, 18 Mar 2014 02:18:40 GMT CRLF
Connection: keep-alive CRLF
Accept-Ranges: bytes CRLF
CRLF
<!doctype html>
<html lang="en">
  <head>
    <meta charset="utf-8">
    <link type="image/x-icon" rel="icon" href="favicon.ico">
    <link type="image/x-icon" rel="shortcut icon" href="favicon.ico">
    <link rel="stylesheet" href="pipe.css">
    ...
```


Request

GET / HTTP/1.1 CRLF

Host: www.nodejs.org CRLF

User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.8; rv:28.0) Gecko/20100101 Firefox/28.0 CRLF

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8 CRLF

Accept-Language: en-us,en;q=0.8,fr;q=0.5,fr-fr;q=0.3 CRLF

Accept-Encoding: gzip, deflate CRLF

Cookie: __utma=212211339.431073283.1392993818.1395308748.1395311696.27;
__utmz=212211339.1395311696.27.19.utmcsr=stackoverflow.com|utmccn=(referral)|utmcmd=referral
|utmcct=/questions/7776452/retrieving-a-list-of-network-interfaces-in-node-js- CRLF

Connection: keep-alive CRLF

CRLF

n header lines

no content (because it is a GET)

1 request line

1 empty line

Response



HTTP/1.1 200 OK CRLF
Server: nginx CRLF
Date: Sat, 05 Apr 2014 11:45:48 GMT CRLF
Content-Type: text/html CRLF
Content-Length: 6368 CRLF
Last-Modified: Tue, 18 Mar 2014 02:18:40 GMT CRLF
Connection: keep-alive CRLF
Accept-Ranges: bytes CRLF
CRLF

n header lines

```
<!doctype html>
<html lang="en">
  <head>
    <meta charset="utf-8">
    <link type="image/x-icon" rel="icon" href="favicon.ico">
    <link type="image/x-icon" rel="shortcut icon" href="favicon.ico">
    <link rel="stylesheet" href="pipe.css">
```

1 status line

1 empty line

body

Resources, Resource Representations & Content Negotiation



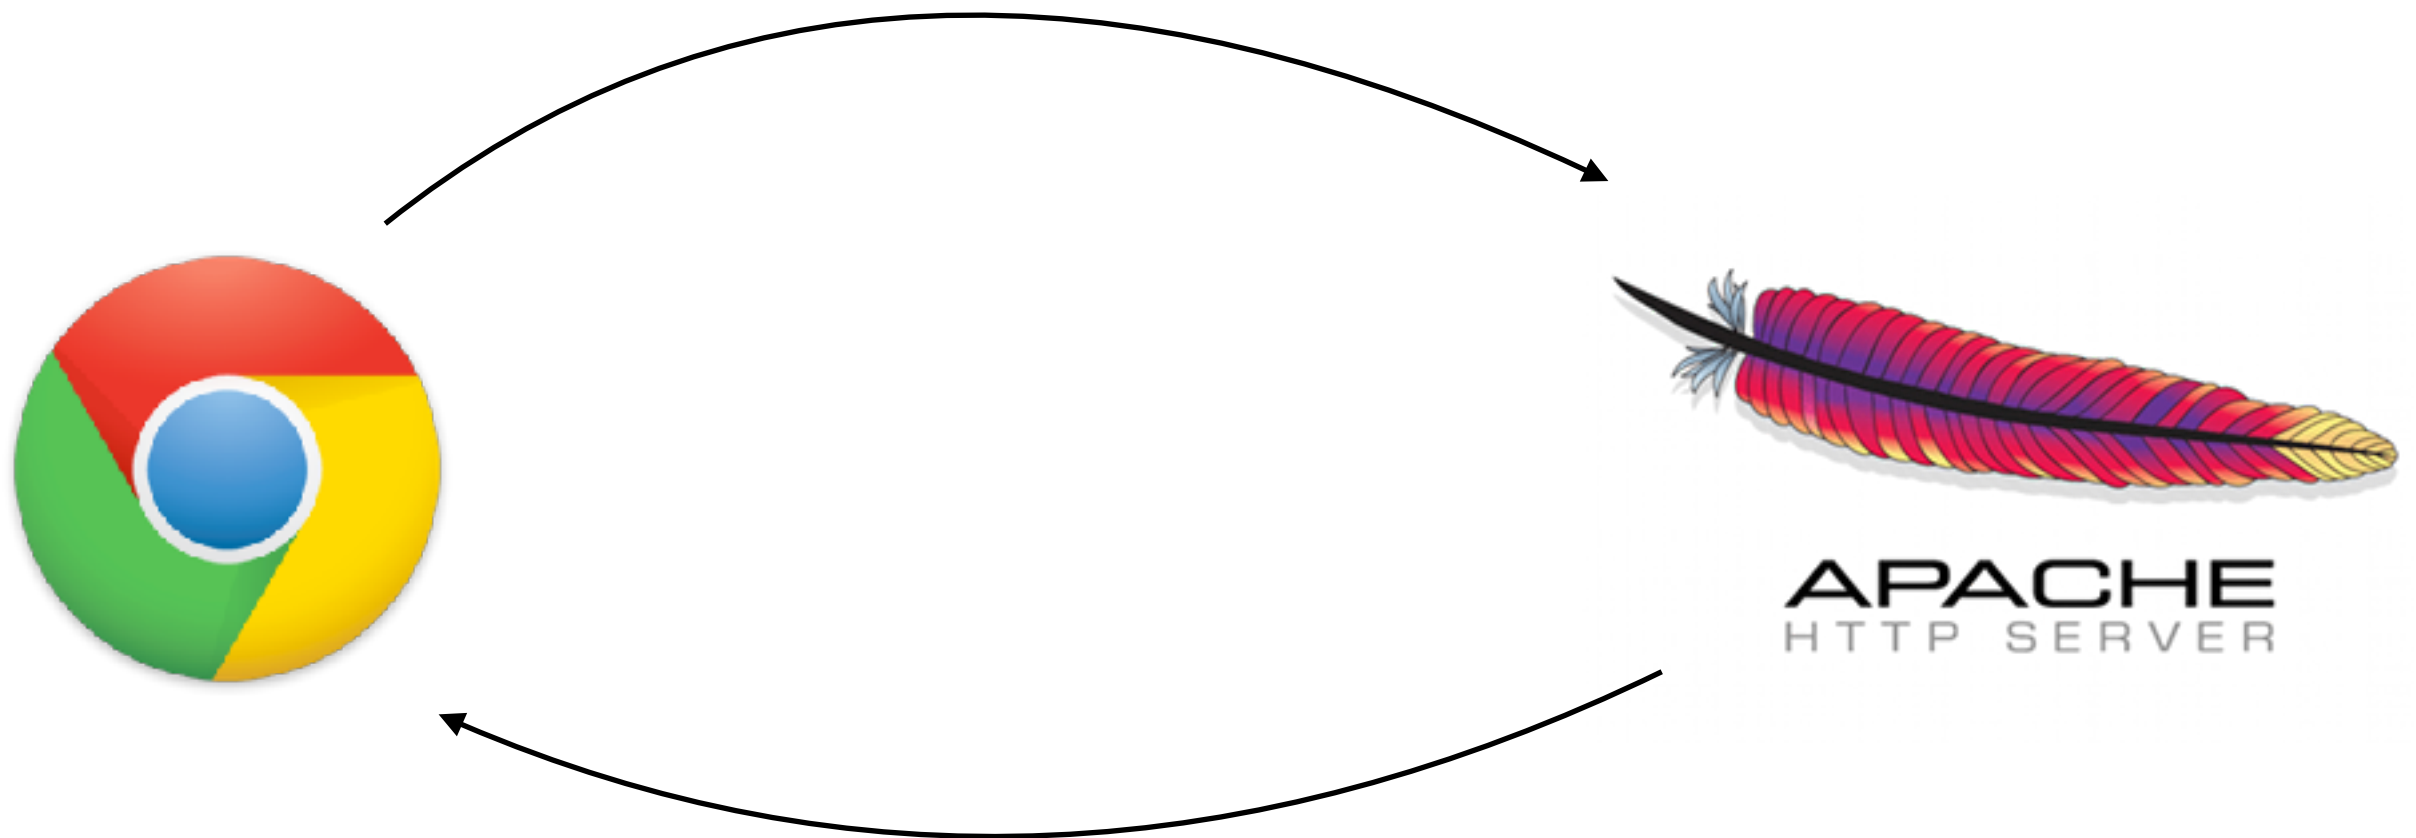
Resource vs Resource Representation

- **The notion of resource is very generic and can represent anything...**
 - An online document
 - A list of online documents
 - A stock quote updated in realtime
 - A vending machine
- **What is transferred is not the resource, but a representation of the resource**
 - HTML representation, JSON representation, PNG representation
 - french representation, english representation, japanese representation
 - etc.

Content Negotiation

- **When making a request, the client specifies its abilities and preferences**
 - media type: image, text, structured text?
 - media format: JSON, XML, etc.?
 - language: english, french, etc.
 - character encoding: UTF-8, ASCII, etc.
- **When answering the request, the server tries to do its best and indicates what it has been able to do**
- **Special headers are used to support this process**
 - **Request:** Accept, Accept-Charset, Accept-Language
 - **Response:** Content-Type, Content-Language

“This is what I am **able to** process, and these are my **preferences...**” (e.g. I am able to deal with plain text and XML, but I prefer JSON)”



“I am **able** to generate XML (not JSON), so you should be able to process this payload”

Protocol Syntax

HTTP Methods

GET

POST

PUT

DELETE

(PATCH)

HTTP Methods & **CRUD** operations

POST

CREATE

GET

READ

PUT

UPDATE

DELETE

DELETE

(PATCH)

PARTIAL UPDATE

URI

<http://www.heig-vd.ch?p1=v1&p2=v2>

query string

Protocol Version

HTTP/1.0

HTTP/1.1

HTTP/2.0

QUIC

HTTP Requests

Full-Request = Request-Line ; Section 5.1
 *(General-Header ; Section 4.3
 | Request-Header ; Section 5.2
 | Entity-Header) ; Section 7.1
 CRLF
 [Entity-Body] ; Section 7.2

Request-Line = Method SP Request-URI SP HTTP-Version CRLF

Request-Header = Authorization ; Section 10.2
 | From ; Section 10.8
 | If-Modified-Since ; Section 10.9
 | Referer ; Section 10.13
 | User-Agent ; Section 10.15

Entity-Header = Allow ; Section 10.1
 | Content-Encoding ; Section 10.3
 | Content-Length ; Section 10.4
 | Content-Type ; Section 10.5
 | Expires ; Section 10.7
 | Last-Modified ; Section 10.10
 | extension-header

HTTP Responses

Full-Response = Status-Line ; Section 6.1
 *(General-Header ; Section 4.3
 | Response-Header ; Section 6.2
 | Entity-Header) ; Section 7.1
 CRLF
 [Entity-Body] ; Section 7.2

Status-Line = HTTP-Version SP Status-Code SP Reason-Phrase CRLF

Response-Header = Location ; Section 10.11
 | Server ; Section 10.14
 | WWW-Authenticate ; Section 10.16

Entity-Header = Allow ; Section 10.1
 | Content-Encoding ; Section 10.3
 | Content-Length ; Section 10.4
 | Content-Type ; Section 10.5
 | Expires ; Section 10.7
 | Last-Modified ; Section 10.10
 | extension-header

Status Codes

The first digit of the Status-Code defines the class of response. The last two digits do not have any categorization role. There are 5 values for the first digit:

- o 1xx: Informational - Not used, but reserved for future use
- o 2xx: Success - The action was successfully received, understood, and accepted.
- o 3xx: Redirection - Further action must be taken in order to complete the request
- o 4xx: Client Error - The request contains bad syntax or cannot be fulfilled
- o 5xx: Server Error - The server failed to fulfill an apparently valid request

Status Codes

Status-Code = "200" ; OK
| "201" ; Created
| "202" ; Accepted
| "204" ; No Content
| "301" ; Moved Permanently
| "302" ; Moved Temporarily
| "304" ; Not Modified
| "400" ; Bad Request
| "401" ; Unauthorized
| "403" ; Forbidden
| "404" ; Not Found
| "500" ; Internal Server Error
| "501" ; Not Implemented
| "502" ; Bad Gateway
| "503" ; Service Unavailable
| extension-code

extension-code = 3DIGIT

Reason-Phrase = *<TEXT, excluding CR, LF>

Parsing HTTP Messages

Process for Parsing HTTP Messages

- **Do NOT read characters, read bytes**
 - At the beginning, you want to parse line by line
 - When consuming the body, you may be dealing with binary content
- **HTTP 1.0**
 - On the client side, read until the connection is closed (end of stream reached).
 - On the server side, use the **Content-Length** header (for POST requests)
- **HTTP 1.1**
 - Static content: use the **Content-Length** header
 - Dynamic content: use the **chunked** transfer encoding

Recommendations

- **Implement your own `LineByLineInputStream`**
 - Remember the lecture about IOs & decorators?
 - You would like to have a `readLine()` method... but this one is available only in `Reader` classes
 - Implement your subclass of `FilterInputStream` and detect `\r\n` sequences
- **Add functionality incrementally, starting with a client**
 - Start with HTTP 1.0 (read until close of connection)
 - Deal with `Content-Length` header
 - Deal with `chunked` transfer encoding

Training

1 Objectifs

Cette semaine, il n'y a pas de laboratoire à rendre. L'objectif est de découvrir le protocole HTTP :

- en utilisant des outils tels que telnet, curl ou postman
- en développant un client HTTP simple (pas besoin de lire toute la spécification)
- en développant un serveur HTTP simple
- en faisant fonctionner le client et le serveur dans 2 containers Docker

Une personne peut travailler sur le client, pendant que l'autre travaille sur le serveur. Une fois la solution complète testée, il est important que les personnes aient bien compris le client et le serveur.

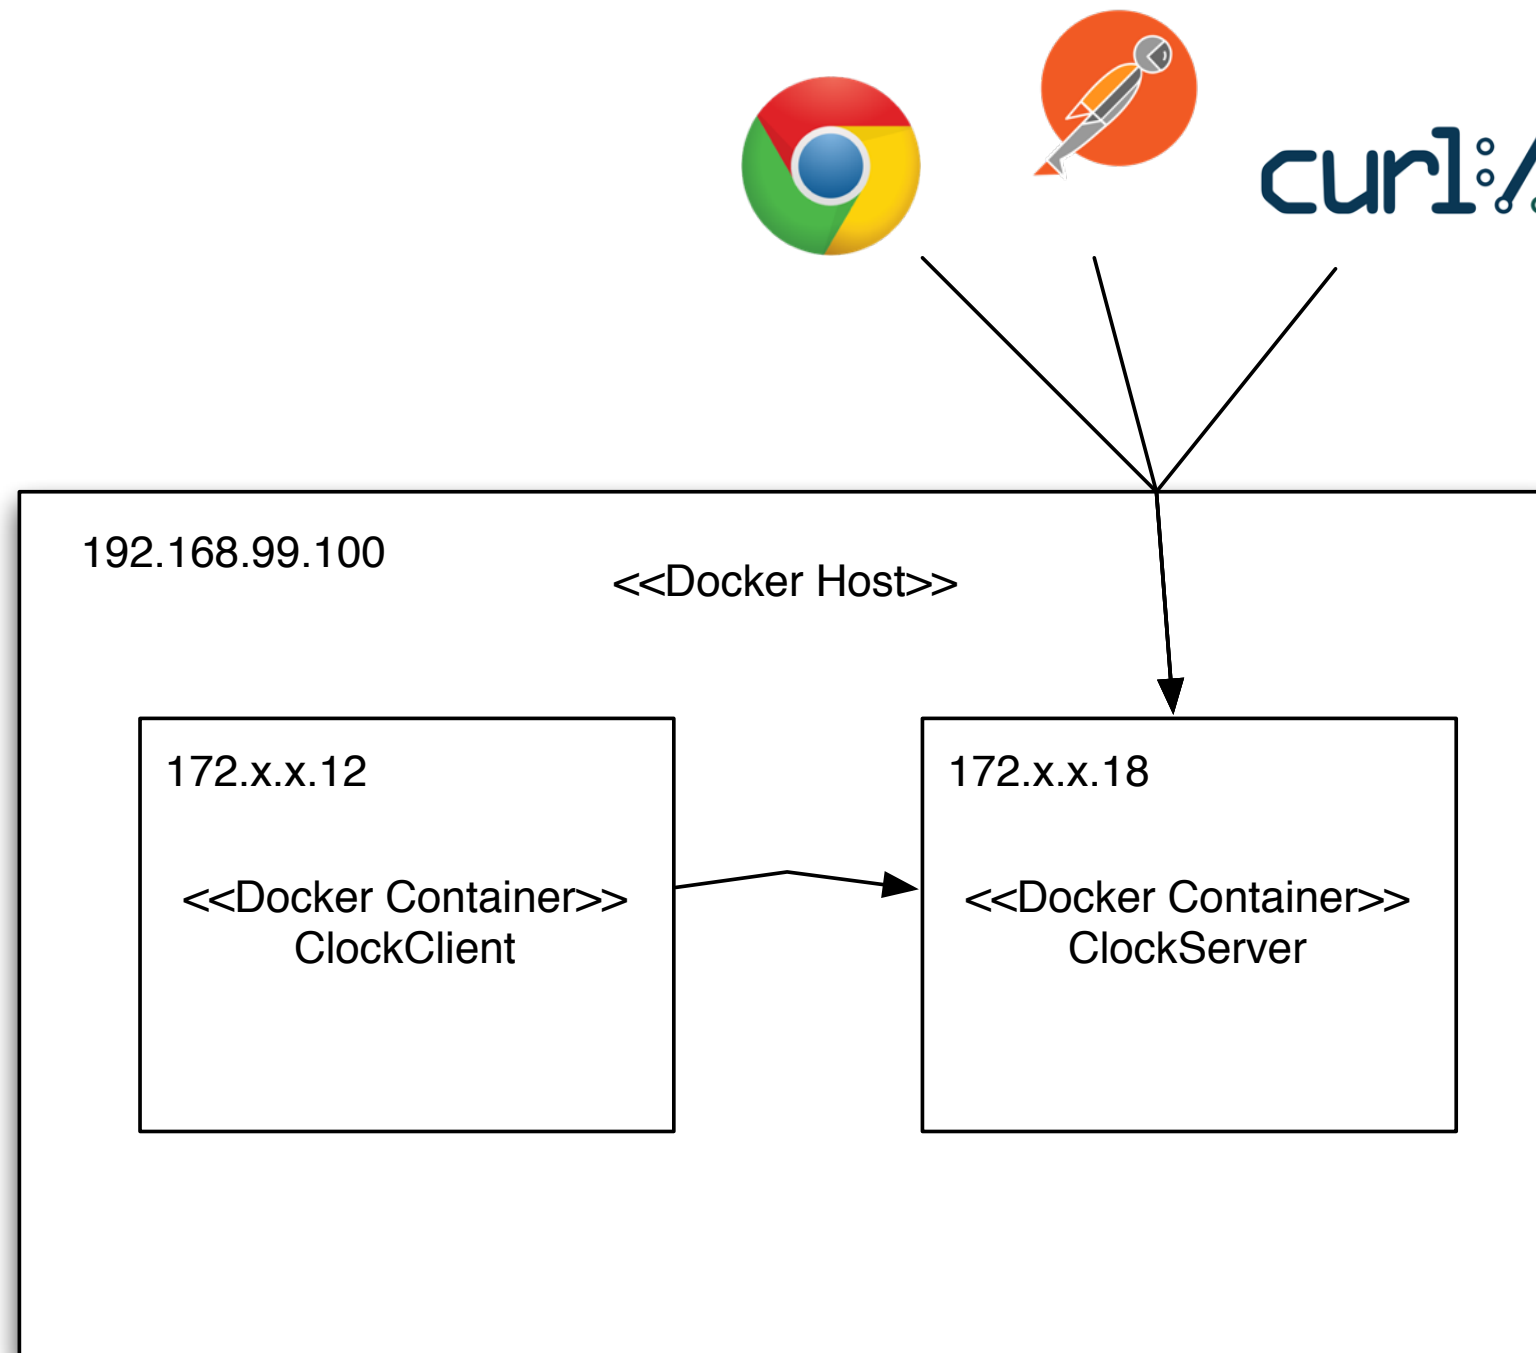
Pour réaliser le client et le serveur, vous pouvez utiliser Java, Javascript ou un autre langage. Vous pouvez décider de travailler avec le niveau d'abstraction qui vous convient (Socket API ou librairie de plus haut niveau). Utiliser la Socket API est ce qui vous permettra comprendre les détails (syntaxe) du protocole HTTP.

2 Conseils

Vous pourriez travailler à deux et réaliser une application avec un client HTTP et un serveur HTTP. Le serveur ne donnerait pas accès à des documents HTML, mais il offrirait un service dynamique.

Par exemple, le serveur pourrait être une horloge et exposer ce service via une interface HTTP :

- Le client peut envoyer des requêtes au serveur pour obtenir l'heure courante (le client fait une lecture avec un GET : il demande au serveur de lui donner des informations). Le serveur doit parser la requête et renvoyer une réponse bien formée. Le client doit la parser et la présenter à l'utilisateur.
- le client peut aussi envoyer des requêtes au serveur pour changer l'heure (le client fait une écriture avec un POST). Le serveur doit lire les données envoyées par le client (dans le body et dans le query string). Il doit répondre au client.
- pour bien comprendre la notion de négociation, faites en sorte que le client puisse spécifier le format qu'il préfère (html, json ou xml). Faites en sorte que le serveur honore ce format quand il prépare la réponse. Dans la même idée, faites en sorte que le client puisse envoyer les données de mise à jour (avec le POST) en json ou en xml.



Useful tools for your
detective arsenal

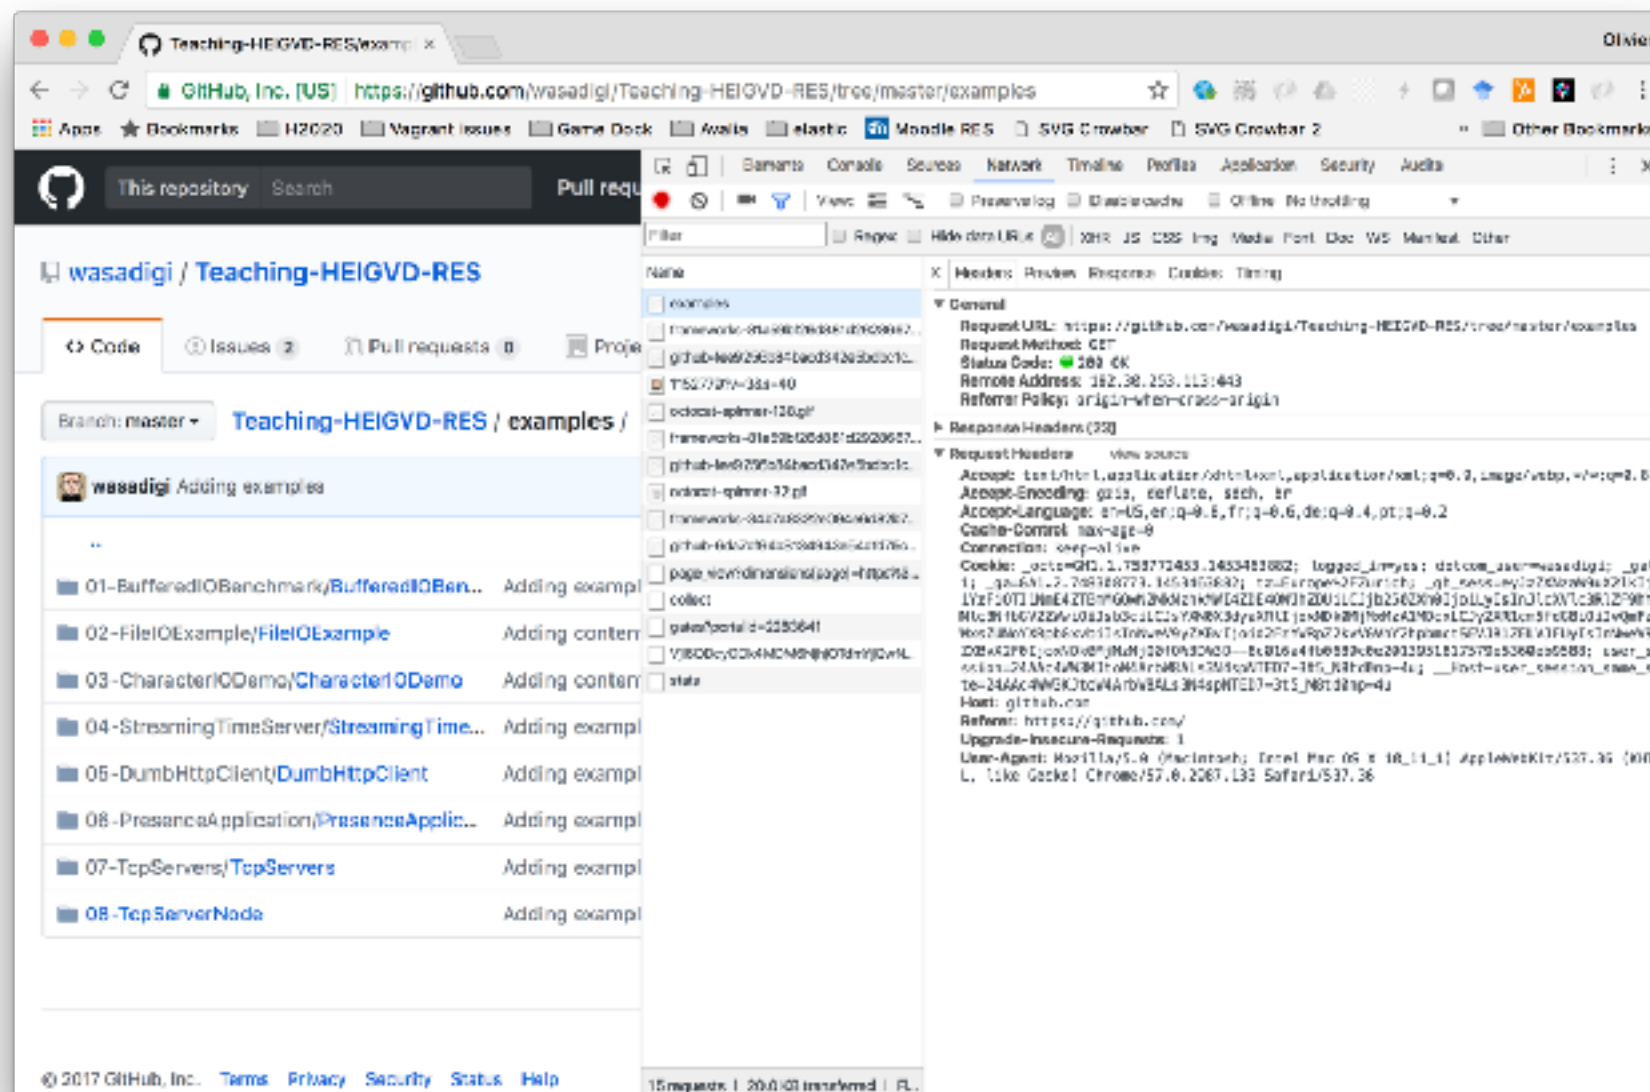


What can I do with it?

inspect request and response headers and body

Where do I find help?

<https://developer.chrome.com/devtools>



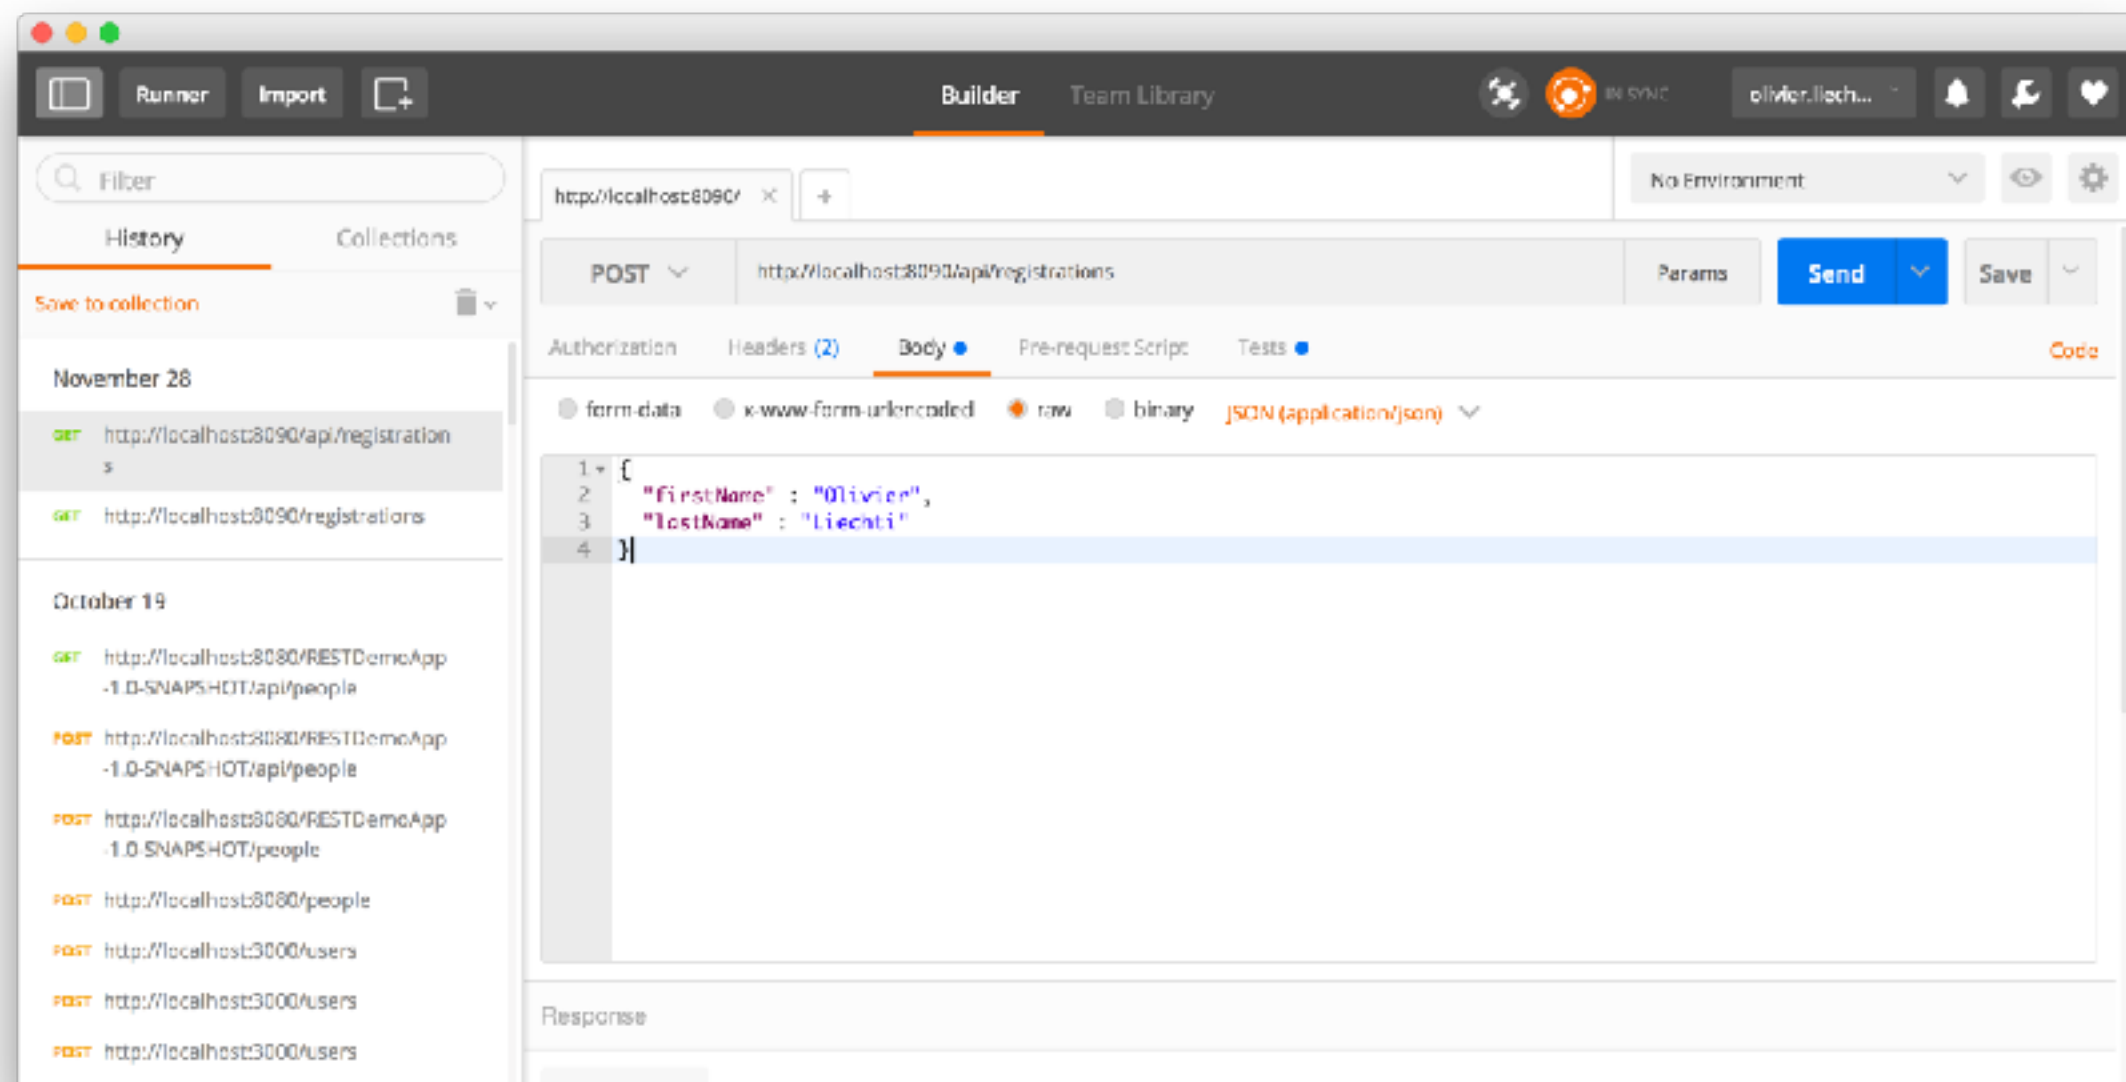


What can I do with it?

Send HTTP requests (also POST, also the body...)

Where do I find help?

<https://www.getpostman.com/>





What can I do with it?	send HTTP requests from terminal, inspect headers
Where do I find help?	https://curl.haxx.se/
How do I install it?	Package manager. You can use it in Docker.



Low level You know the drill, just implement the protocol syntax!	<pre>Socket s = new Socket("www.heig-vd.ch", 80); ...</pre>
Intermediate level (see Java tutorial)	<pre>URL url = new URL("www.heig-vd.ch"); URLConnection con = url.openConnection();</pre>
High level (see http://hc.apache.org/ and don't forget to add the dependency in your pom.xml)	<pre>Request.Get("http://somehost/") .connectTimeout(1000) .socketTimeout(1000) .execute().returnContent().asString();</pre>



Low level	<pre>var client = new net.Socket(); client.connect(80, 'www.heig-vd.ch', function() {...}); client.on('data', function(data) {...});</pre>
Intermediate level (see Java tutorial)	<pre>var http = require('http'); var options = {...}; var request = http.request(options, function(response) {...}); request.end();</pre>
High level (many npm modules: request, request-promise, superagent)	<pre>request .post('/api/pet') .send({ name: 'Manny', species: 'cat' }) .set('Accept', 'application/json') .end(function(err, res){ ... });</pre>

The HTTP Host Header

In the old days...

DNS

`www.company-a.ch` → public IP address
`193.134.218.45`

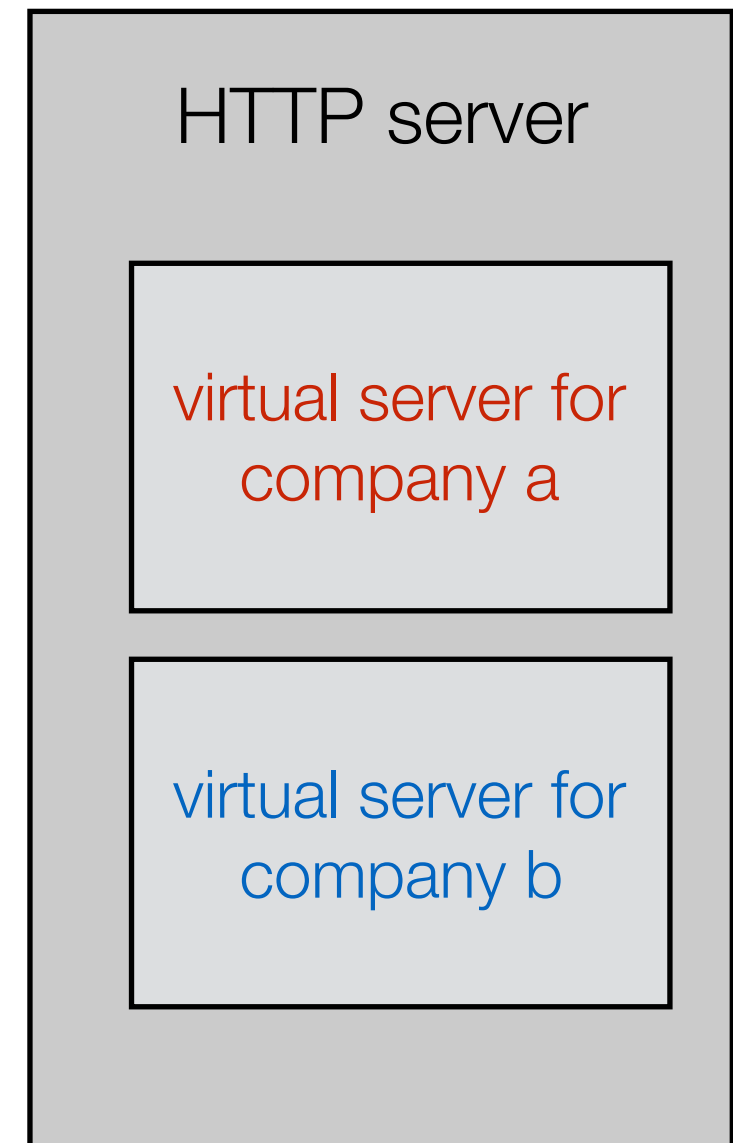
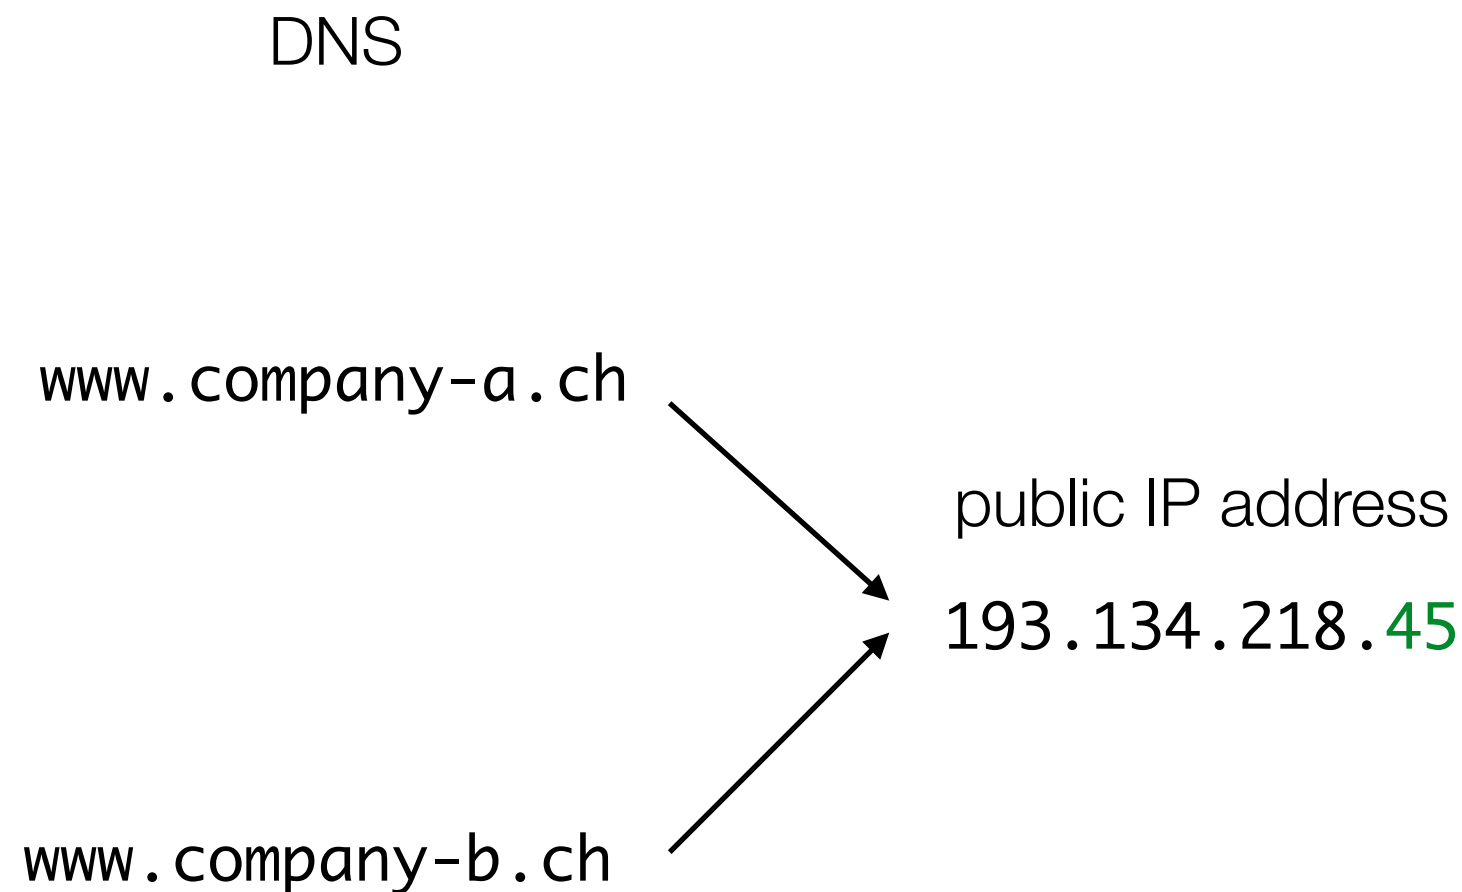
HTTP server
with content for
company a

`www.company-b.ch` → public IP address
`193.134.218.46`

HTTP server
with content for
company b

We would run out of IP addresses! A big web hosting provider would need thousands of public IPs.

Virtual hosting



Ok... but when the HTTP server receives a request, how does it know if the user wants content A or content B?

HTTP 1.0

GET /index.html HTTP/1.0
Accept: text/html

With this information, impossible to know if the user wants to access virtual host A or B. We can only have one web site per IP address

HTTP 1.1

GET /index.html HTTP/1.0
Host: **www.company-a.ch**
Accept: text/html

GET /index.html HTTP/1.0
Host: **www.company-b.ch**
Accept: text/html

With HTTP 1.1, the Host header is mandatory and is used to solve this problem.
The HTTP server can “forward” the request to the relevant virtual host

HTTP is a Stateless
Request-Reply Transfer Protocol

Stateless Protocol... but Stateful Applications!



Managing State on Top of HTTP

- **Approach 1: move the state back-and-forth**
 - One way to do it is to use **hidden fields** in HTML forms
- **Approach 2 (better): maintain state on the backend, transfer session IDs**
 - One way to do it is to use parameters in the **query string** (security...)
 - One way is to use **cookies**



Passing State Back and Forth

C: Hello, I am new here. My name is Bob.

S: Welcome, let's have a chat [You told me that "My name is Bob"].

C: [My name is Bob]. What's the time?

S: Hi again Bob. It's 10:45 AM. [You told me that "My name is Bob". You asked me what is the time]

#	Name	Description
1	Mark	Otto
2	Jacob	Thornton

Passing Session ID Back and Forth

C: Hello, I am new here. My name is Bob.

S: Welcome Bob, let's have a chat. Your session id is 42.

C: My session id is 42. What's the time?

S: -- checking my notes... hum... ok, I found what I remember about session 42...

S: Hi again Bob. It's 10:45 AM.

C: My session id is 42. How do you do?

S: -- checking my notes... hum... ok, I found what I remember about session 42...

S: I am fine, Bob, thank you.

C: My session id is 42. How do you do?

S: -- checking my notes... hum... ok, I found what I remember about session 42...

S: I told you I am fine... are you stupid or what?

C: My session id is 42. If you take it like that, I am gone.
Forever.

S: -- checking my notes... hum... ok, I found what I remember about session 42...

S: -- putting 42 file into trash...

S: Bye Bob.

Passing Session ID in URL

GET /login HTTP/1.1
Host: intra.heig-vd.ch

HTTP/1.1 302 Found
Location: <http://intra.heig-vd.ch/home?sessionId=83939>

GET /home?sessionId=83939 HTTP/1.1
Host: intra.heig-vd.ch

GET /page2?sessionId=83939 HTTP/1.1
Host: intra.heig-vd.ch

GET /logout?sessionId=83939 HTTP/1.1
Host: intra.heig-vd.ch

Passing Session ID in Cookies

- The client **sends a request to the server for the first time**. It cannot send any cookie (it does not have any!).
- The server **understands that it is a new client**. It creates a session and generates a unique session ID.
- In the HTTP response, the server sends the unique session ID in a **Set-Cookie header** (“Here is a token... next time you come to see me, show me the token and I will recognize you!”).
- The **client keeps track of the cookie** either in memory or on disk (it has a cookie store, i.e. a list of cookies sent by servers).
- Before the client sends a second request to the server, it lookups the cookie store: “I am about to visit a server... did I receive any cookie from it?”. If yes, it sends the cookie in a **Cookie header**.

Passing Session ID in Cookies

- The server sends a cookie (token):

Set-Cookie: GAPSSESSID=4rm?????na; path=/; domain=gaps.heig-vd.ch

- The client sends it each time it returns to the server:

Cookie: GAPSSESSID=4rm?????na

Where does the server store the state?

- What is transferred in the cookie is **only a session ID**, not session data (and BTW not user credentials!).
- The **data associated to the session** typically includes the user identity (and roles), the previous actions performed by the user, the preferences, etc.
- **This data can be stored in various places.** Where you store it can have a big impact on performance, scalability and fault-tolerance. There is no universal “best solution” - the choice depends on many factors.
- Choice 1: **store everything in the DB.** One advantage is that if a web server crashes, it does not bring down the state. One disadvantage is that performance may
- Choice 2: **keep in the RAM of the web server.** One advantage is the performance. One disadvantage is that you have to use sticky sessions. Another disadvantage is that RAM consumption could be high.
- Choice 3: **use a distributed caching layer** (e.g. redis, memcached). One advantage is fault tolerance and scalability. One disadvantage is higher complexity and cost.